```
t/session.py", line 950, in run
    run_metadata_ptr)
  File "/home/ubuntu/resnet/venv/lib/python3.6/site-packages/tensorflow/python/clien
t/session.py", line 1173, in _run
    feed_dict_tensor, options, run_metadata)
  File "/home/ubuntu/resnet/venv/lib/python3.6/site-packages/tensorflow/python/clien
t/session.py", line 1350, in _do_run
    run_metadata)
  File "/home/ubuntu/resnet/venv/lib/python3.6/site-packages/tensorflow/python/clien
t/session.py", line 1356, in _do_call
    return fn(*args)
  File "/home/ubuntu/resnet/venv/lib/python3.6/site-packages/tensorflow/python/clien
t/session.py", line 1341, in _run_fn
    options, feed_dict, fetch_list, target_list, run_metadata)
  File "/home/ubuntu/resnet/venv/lib/python3.6/site-packages/tensorflow/python/clien
t/session.py", line 1429, in _call_tf_sessionrun
    run_metadata)
KeyboardInterrupt
(venv) ubuntu@ip-172-31-88-204:~/resnet$
(venv) ubuntu@ip-172-31-88-204:~/resnet$ ls
README.md    config.yaml  logs     requirements.txt  rotnet.py
__pycache__  data.py      main.py  resnet.py         venv
(venv) ubuntu@ip-172-31-88-204:~/resnet$ clear

(venv) ubuntu@ip-172-31-88-204:~/resnet$ ls
README.md    config.yaml  logs     requirements.txt  rotnet.py
__pycache__  data.py      main.py  resnet.py         venv
(venv) ubuntu@ip-172-31-88-204:~/resnet$ vim config.yaml
(venv) ubuntu@ip-172-31-88-204:~/resnet$ python3 main.py --config config.yaml --trai
n --data_dir ./data/cifar-10-batches-py/ --model_number 1
2020-04-13 07:18:38.614468: I tensorflow/core/platform/cpu_feature_guard.cc:142] You
r CPU supports instructions that this TensorFlow binary was not compiled to use: AVX
2 FMA
2020-04-13 07:18:38.636385: I tensorflow/core/platform/profile_utils/cpu_utils.cc:94
] CPU Frequency: 2300050000 Hz
2020-04-13 07:18:38.636791: I tensorflow/compiler/xla/service/service.cc:168] XLA se
rvice 0x4332da0 executing computations on platform Host. Devices:
2020-04-13 07:18:38.636833: I tensorflow/compiler/xla/service/service.cc:175]   Stre
amExecutor device (0): <undefined>, <undefined>
[INFO] Reading configuration file
WARNING: Logging before flag parsing goes to stderr.
W0413 07:18:39.030417 139687167268672 lazy_loader.py:50]
The TensorFlow contrib module will not be included in TensorFlow 2.0.
For more information, please see:
  * https://github.com/tensorflow/community/blob/master/rfcs/20180907-contrib-sunset
.md
  * https://github.com/tensorflow/addons
  * https://github.com/tensorflow/io (for I/O related ops)
If you depend on functionality not listed there, please file an issue.

(?, 1, 1, 512)
W0413 07:18:40.508047 139687167268672 deprecation.py:506] From /home/ubuntu/resnet/v
env/lib/python3.6/site-packages/tensorflow/python/ops/init_ops.py:1251: calling Vari
anceScaling.__init__ (from tensorflow.python.ops.init_ops) with dtype is deprecated
and will be removed in a future version.
```

```
Instructions for updating:
Call initializer instance with the dtype argument instead of passing it to the const
ructor
[name: "/device:CPU:0"
device_type: "CPU"
memory_limit: 268435456
locality {
}
incarnation: 16844950109245577651
, name: "/device:XLA_CPU:0"
device_type: "XLA_CPU"
memory_limit: 17179869184
locality {
}
incarnation: 6515080501586048891
physical_device_desc: "device: XLA_CPU device"
]
2020-04-13 07:18:44.492984: W tensorflow/compiler/jit/mark_for_compilation_pass.cc:1
412] (One-time warning): Not using XLA:CPU for cluster because envvar TF_XLA_FLAGS=-
-tf_xla_cpu_global_jit was not set.  If you want XLA:CPU, either set that envvar, or
 use experimental_jit_scope to enable XLA:CPU.  To confirm that XLA is active, pass
--vmodule=xla_compilation_cache=1 (as a proper command-line flag, not via TF_XLA_FLA
GS) or set the envvar XLA_FLAGS=--xla_hlo_profile.
[INFO] Getting Training Data
X_train shape:
(134000, 32, 32, 3)
y_train shape:
(134000, 4)
X_val shape:
(66000, 32, 32, 3)
y_val shape:
(66000, 4)
[INFO] Starting Training...
Epoch: 0, Batch: 0 ==> Accuracy: 0.220703125, Loss: 20.171598434448242
Epoch: 0, Batch: 1 ==> Accuracy: 0.240234375, Loss: 94.45580291748047
Epoch: 0, Batch: 2 ==> Accuracy: 0.25, Loss: 44.4500732421875
Epoch: 0, Batch: 3 ==> Accuracy: 0.234375, Loss: 25.50696563720703
Epoch: 0, Batch: 4 ==> Accuracy: 0.251953125, Loss: 8.772745132446289
Epoch: 0, Batch: 5 ==> Accuracy: 0.255859375, Loss: 2.8025922775268555
Epoch: 0, Batch: 6 ==> Accuracy: 0.26953125, Loss: 1.8570579290390015
Epoch: 0, Batch: 7 ==> Accuracy: 0.24609375, Loss: 2.010237693786621
Epoch: 0, Batch: 8 ==> Accuracy: 0.255859375, Loss: 1.4459054470062256
Epoch: 0, Batch: 9 ==> Accuracy: 0.279296875, Loss: 1.5592052936553955
Epoch: 0, Batch: 10 ==> Accuracy: 0.2734375, Loss: 1.5494649410247803
Epoch: 0, Batch: 11 ==> Accuracy: 0.23828125, Loss: 1.4510706663131714
Epoch: 0, Batch: 12 ==> Accuracy: 0.25390625, Loss: 1.4100409746170044
Epoch: 0, Batch: 13 ==> Accuracy: 0.28515625, Loss: 1.4462618827819824
Epoch: 0, Batch: 14 ==> Accuracy: 0.248046875, Loss: 1.4243677854537964
Epoch: 0, Batch: 15 ==> Accuracy: 0.23828125, Loss: 1.3925795555114746
Epoch: 0, Batch: 16 ==> Accuracy: 0.2734375, Loss: 1.3722972869873047
Epoch: 0, Batch: 17 ==> Accuracy: 0.220703125, Loss: 1.4323196411132812
Epoch: 0, Batch: 18 ==> Accuracy: 0.251953125, Loss: 1.4157569408416748
Epoch: 0, Batch: 19 ==> Accuracy: 0.30859375, Loss: 1.3715534210205078
Epoch: 0, Batch: 20 ==> Accuracy: 0.359375, Loss: 1.3533873558044434
Epoch: 0, Batch: 21 ==> Accuracy: 0.318359375, Loss: 1.363830804824829
```

```
Epoch: 0, Batch: 22 ==> Accuracy: 0.380859375, Loss: 1.341454267501831
Epoch: 0, Batch: 23 ==> Accuracy: 0.333984375, Loss: 1.3711951971054077
Epoch: 0, Batch: 24 ==> Accuracy: 0.326171875, Loss: 1.3588218688964844
Epoch: 0, Batch: 25 ==> Accuracy: 0.345703125, Loss: 1.3542966842651367
Epoch: 0, Batch: 26 ==> Accuracy: 0.3125, Loss: 1.3511087894439697
Epoch: 0, Batch: 27 ==> Accuracy: 0.36328125, Loss: 1.3496513366699219
Epoch: 0, Batch: 28 ==> Accuracy: 0.328125, Loss: 1.3482518196105957
Epoch: 0, Batch: 29 ==> Accuracy: 0.392578125, Loss: 1.3184359073638916
Epoch: 0, Batch: 30 ==> Accuracy: 0.333984375, Loss: 1.3497532606124878
Epoch: 0, Batch: 31 ==> Accuracy: 0.37109375, Loss: 1.3180487155914307
Epoch: 0, Batch: 32 ==> Accuracy: 0.35546875, Loss: 1.3298789262771606
Epoch: 0, Batch: 33 ==> Accuracy: 0.447265625, Loss: 1.2807639837265015
Epoch: 0, Batch: 34 ==> Accuracy: 0.421875, Loss: 1.308258056640625
Epoch: 0, Batch: 35 ==> Accuracy: 0.470703125, Loss: 1.2541190385818481
Epoch: 0, Batch: 36 ==> Accuracy: 0.435546875, Loss: 1.2648732662200928
Epoch: 0, Batch: 37 ==> Accuracy: 0.408203125, Loss: 1.3057856559753418
Epoch: 0, Batch: 38 ==> Accuracy: 0.416015625, Loss: 1.297888994216919
Epoch: 0, Batch: 39 ==> Accuracy: 0.44921875, Loss: 1.2579550743103027
Epoch: 0, Batch: 40 ==> Accuracy: 0.443359375, Loss: 1.2636964321136475
Epoch: 0, Batch: 41 ==> Accuracy: 0.400390625, Loss: 1.290907859802246
Epoch: 0, Batch: 42 ==> Accuracy: 0.423828125, Loss: 1.247591257095337
Epoch: 0, Batch: 43 ==> Accuracy: 0.41796875, Loss: 1.2864670753479004
Epoch: 0, Batch: 44 ==> Accuracy: 0.4296875, Loss: 1.267054557800293
Epoch: 0, Batch: 45 ==> Accuracy: 0.423828125, Loss: 1.2695443630218506
Epoch: 0, Batch: 46 ==> Accuracy: 0.4296875, Loss: 1.275054693222046
Epoch: 0, Batch: 47 ==> Accuracy: 0.453125, Loss: 1.2336311340332031
Epoch: 0, Batch: 48 ==> Accuracy: 0.44140625, Loss: 1.2378356456756592
Epoch: 0, Batch: 49 ==> Accuracy: 0.451171875, Loss: 1.246256947517395
Epoch: 0, Batch: 50 ==> Accuracy: 0.41796875, Loss: 1.2350292205810547
Epoch: 0, Batch: 51 ==> Accuracy: 0.47265625, Loss: 1.194960594177246
Epoch: 0, Batch: 52 ==> Accuracy: 0.44140625, Loss: 1.2735174894332886
Epoch: 0, Batch: 53 ==> Accuracy: 0.44921875, Loss: 1.2064578533172607
Epoch: 0, Batch: 54 ==> Accuracy: 0.458984375, Loss: 1.177692174911499
Epoch: 0, Batch: 55 ==> Accuracy: 0.455078125, Loss: 1.2291741371154785
Epoch: 0, Batch: 56 ==> Accuracy: 0.490234375, Loss: 1.1775593757629395
Epoch: 0, Batch: 57 ==> Accuracy: 0.478515625, Loss: 1.1585958003997803
Epoch: 0, Batch: 58 ==> Accuracy: 0.50390625, Loss: 1.1394697427749634
Epoch: 0, Batch: 59 ==> Accuracy: 0.431640625, Loss: 1.2361798286437988
Epoch: 0, Batch: 60 ==> Accuracy: 0.42578125, Loss: 1.2161043882369995
Epoch: 0, Batch: 61 ==> Accuracy: 0.45703125, Loss: 1.2038389444351196
Epoch: 0, Batch: 62 ==> Accuracy: 0.443359375, Loss: 1.1936452388763428
Epoch: 0, Batch: 63 ==> Accuracy: 0.5, Loss: 1.1799407005310059
Epoch: 0, Batch: 64 ==> Accuracy: 0.462890625, Loss: 1.167911171913147
Epoch: 0, Batch: 65 ==> Accuracy: 0.4453125, Loss: 1.2044072151184082
Epoch: 0, Batch: 66 ==> Accuracy: 0.498046875, Loss: 1.1718876361846924
Epoch: 0, Batch: 67 ==> Accuracy: 0.41796875, Loss: 1.2252217531204224
Epoch: 0, Batch: 68 ==> Accuracy: 0.4765625, Loss: 1.1508368253707886
Epoch: 0, Batch: 69 ==> Accuracy: 0.482421875, Loss: 1.1658200025558472
Epoch: 0, Batch: 70 ==> Accuracy: 0.458984375, Loss: 1.1780855655670166
Epoch: 0, Batch: 71 ==> Accuracy: 0.45703125, Loss: 1.1483519077301025
Epoch: 0, Batch: 72 ==> Accuracy: 0.50390625, Loss: 1.122720718383789
Epoch: 0, Batch: 73 ==> Accuracy: 0.47265625, Loss: 1.1833548545837402
Epoch: 0, Batch: 74 ==> Accuracy: 0.482421875, Loss: 1.1724894046783447
Epoch: 0, Batch: 75 ==> Accuracy: 0.5, Loss: 1.1239798069000244
Epoch: 0, Batch: 76 ==> Accuracy: 0.490234375, Loss: 1.1389987468719482
```

```
Epoch: 0, Batch: 77 ==> Accuracy: 0.501953125, Loss: 1.1233370304107666
Epoch: 0, Batch: 78 ==> Accuracy: 0.48828125, Loss: 1.1360210180282593
Epoch: 0, Batch: 79 ==> Accuracy: 0.52734375, Loss: 1.0961554050445557
Epoch: 0, Batch: 80 ==> Accuracy: 0.525390625, Loss: 1.063328742980957
Epoch: 0, Batch: 81 ==> Accuracy: 0.494140625, Loss: 1.1129517555236816
Epoch: 0, Batch: 82 ==> Accuracy: 0.5078125, Loss: 1.1440014839172363
Epoch: 0, Batch: 83 ==> Accuracy: 0.525390625, Loss: 1.1080586910247803
Epoch: 0, Batch: 84 ==> Accuracy: 0.546875, Loss: 1.057712435722351
Epoch: 0, Batch: 85 ==> Accuracy: 0.48828125, Loss: 1.1284152269363403
Epoch: 0, Batch: 86 ==> Accuracy: 0.5078125, Loss: 1.13847017288208
Epoch: 0, Batch: 87 ==> Accuracy: 0.5, Loss: 1.114682912826538
Epoch: 0, Batch: 88 ==> Accuracy: 0.46484375, Loss: 1.1474958658218384
Epoch: 0, Batch: 89 ==> Accuracy: 0.5390625, Loss: 1.1032934188842773
Epoch: 0, Batch: 90 ==> Accuracy: 0.51953125, Loss: 1.0913411378860474
Epoch: 0, Batch: 91 ==> Accuracy: 0.46875, Loss: 1.1399214267730713
Epoch: 0, Batch: 92 ==> Accuracy: 0.509765625, Loss: 1.104689121246338
Epoch: 0, Batch: 93 ==> Accuracy: 0.525390625, Loss: 1.095777988433838
Epoch: 0, Batch: 94 ==> Accuracy: 0.501953125, Loss: 1.1341843605041504
Epoch: 0, Batch: 95 ==> Accuracy: 0.5078125, Loss: 1.1326054334640503
Epoch: 0, Batch: 96 ==> Accuracy: 0.501953125, Loss: 1.169696569442749
Epoch: 0, Batch: 97 ==> Accuracy: 0.509765625, Loss: 1.0902959108352661
Epoch: 0, Batch: 98 ==> Accuracy: 0.47265625, Loss: 1.151378870010376
Epoch: 0, Batch: 99 ==> Accuracy: 0.533203125, Loss: 1.074049711227417
Epoch: 0, Batch: 100 ==> Accuracy: 0.486328125, Loss: 1.1140303611755371
Epoch: 0, Batch: 101 ==> Accuracy: 0.513671875, Loss: 1.0830347537994385
Epoch: 0, Batch: 102 ==> Accuracy: 0.4765625, Loss: 1.122849702835083
Epoch: 0, Batch: 103 ==> Accuracy: 0.5234375, Loss: 1.076697826385498
Epoch: 0, Batch: 104 ==> Accuracy: 0.5, Loss: 1.1088316440582275
Epoch: 0, Batch: 105 ==> Accuracy: 0.55078125, Loss: 1.075404405593872
Epoch: 0, Batch: 106 ==> Accuracy: 0.513671875, Loss: 1.0820391178131104
Epoch: 0, Batch: 107 ==> Accuracy: 0.5078125, Loss: 1.0982670783996582
Epoch: 0, Batch: 108 ==> Accuracy: 0.50390625, Loss: 1.1111501455307007
Epoch: 0, Batch: 109 ==> Accuracy: 0.525390625, Loss: 1.0852999687194824
Epoch: 0, Batch: 110 ==> Accuracy: 0.486328125, Loss: 1.1602033376693726
Epoch: 0, Batch: 111 ==> Accuracy: 0.53125, Loss: 1.0934966802597046
Epoch: 0, Batch: 112 ==> Accuracy: 0.505859375, Loss: 1.06764817237854
Epoch: 0, Batch: 113 ==> Accuracy: 0.49609375, Loss: 1.1357479095458984
Epoch: 0, Batch: 114 ==> Accuracy: 0.552734375, Loss: 1.0618480443954468
Epoch: 0, Batch: 115 ==> Accuracy: 0.56640625, Loss: 1.025792121887207
Epoch: 0, Batch: 116 ==> Accuracy: 0.505859375, Loss: 1.1193177700042725
Epoch: 0, Batch: 117 ==> Accuracy: 0.560546875, Loss: 1.0687532424926758
Epoch: 0, Batch: 118 ==> Accuracy: 0.54296875, Loss: 1.0716947317123413
Epoch: 0, Batch: 119 ==> Accuracy: 0.53515625, Loss: 1.0847892761230469
Epoch: 0, Batch: 120 ==> Accuracy: 0.564453125, Loss: 1.0369740724563599
Epoch: 0, Batch: 121 ==> Accuracy: 0.548828125, Loss: 1.077918529510498
Epoch: 0, Batch: 122 ==> Accuracy: 0.525390625, Loss: 1.099522352218628
Epoch: 0, Batch: 123 ==> Accuracy: 0.53125, Loss: 1.0234830379486084
Epoch: 0, Batch: 124 ==> Accuracy: 0.5546875, Loss: 1.0336389541625977
Epoch: 0, Batch: 125 ==> Accuracy: 0.52734375, Loss: 1.0623536109924316
Epoch: 0, Batch: 126 ==> Accuracy: 0.501953125, Loss: 1.137070655822754
Epoch: 0, Batch: 127 ==> Accuracy: 0.53515625, Loss: 1.0785822868347168
Epoch: 0, Batch: 128 ==> Accuracy: 0.537109375, Loss: 1.0703089237213135
Epoch: 0, Batch: 129 ==> Accuracy: 0.525390625, Loss: 1.0642340183258057
Epoch: 0, Batch: 130 ==> Accuracy: 0.5234375, Loss: 1.085536003112793
Epoch: 0, Batch: 131 ==> Accuracy: 0.564453125, Loss: 1.0093945264816284
```

```
Epoch: 0, Batch: 132 ==> Accuracy: 0.564453125, Loss: 1.0329053401947021
Epoch: 0, Batch: 133 ==> Accuracy: 0.525390625, Loss: 1.0538299083709717
Epoch: 0, Batch: 134 ==> Accuracy: 0.576171875, Loss: 1.030556559562683
Epoch: 0, Batch: 135 ==> Accuracy: 0.51953125, Loss: 1.1137665510177612
Epoch: 0, Batch: 136 ==> Accuracy: 0.587890625, Loss: 1.021773338317871
Epoch: 0, Batch: 137 ==> Accuracy: 0.55859375, Loss: 1.0414669513702393
Epoch: 0, Batch: 138 ==> Accuracy: 0.560546875, Loss: 1.0618128776550293
Epoch: 0, Batch: 139 ==> Accuracy: 0.564453125, Loss: 1.0489494800567627
Epoch: 0, Batch: 140 ==> Accuracy: 0.52734375, Loss: 1.0666191577911377
Epoch: 0, Batch: 141 ==> Accuracy: 0.51171875, Loss: 1.083989143371582
Epoch: 0, Batch: 142 ==> Accuracy: 0.54296875, Loss: 1.045447587966919
Epoch: 0, Batch: 143 ==> Accuracy: 0.498046875, Loss: 1.109662413597107
Epoch: 0, Batch: 144 ==> Accuracy: 0.59375, Loss: 0.9716939926147461
Epoch: 0, Batch: 145 ==> Accuracy: 0.564453125, Loss: 1.0494177341461182
Epoch: 0, Batch: 146 ==> Accuracy: 0.568359375, Loss: 1.0436975955963135
Epoch: 0, Batch: 147 ==> Accuracy: 0.54296875, Loss: 1.1048541069030762
Epoch: 0, Batch: 148 ==> Accuracy: 0.552734375, Loss: 1.0020909309387207
Epoch: 0, Batch: 149 ==> Accuracy: 0.564453125, Loss: 1.0438597202301025
Epoch: 0, Batch: 150 ==> Accuracy: 0.541015625, Loss: 1.0628101825714111
Epoch: 0, Batch: 151 ==> Accuracy: 0.591796875, Loss: 1.0327062606811523
Epoch: 0, Batch: 152 ==> Accuracy: 0.54296875, Loss: 1.087348461151123
Epoch: 0, Batch: 153 ==> Accuracy: 0.580078125, Loss: 1.0265153646469116
Epoch: 0, Batch: 154 ==> Accuracy: 0.6015625, Loss: 0.9842114448547363
Epoch: 0, Batch: 155 ==> Accuracy: 0.568359375, Loss: 1.0673410892486572
Epoch: 0, Batch: 156 ==> Accuracy: 0.560546875, Loss: 1.0385973453521729
Epoch: 0, Batch: 157 ==> Accuracy: 0.55078125, Loss: 1.0175034999847412
Epoch: 0, Batch: 158 ==> Accuracy: 0.576171875, Loss: 1.0136961936950684
Epoch: 0, Batch: 159 ==> Accuracy: 0.55859375, Loss: 1.0581692457199097
Epoch: 0, Batch: 160 ==> Accuracy: 0.564453125, Loss: 1.0180598497390747
Epoch: 0, Batch: 161 ==> Accuracy: 0.560546875, Loss: 1.0338033437728882
Epoch: 0, Batch: 162 ==> Accuracy: 0.515625, Loss: 1.0723239183425903
Epoch: 0, Batch: 163 ==> Accuracy: 0.548828125, Loss: 1.0447226762771606
Epoch: 0, Batch: 164 ==> Accuracy: 0.525390625, Loss: 1.0808672904968262
Epoch: 0, Batch: 165 ==> Accuracy: 0.564453125, Loss: 0.9874070286750793
Epoch: 0, Batch: 166 ==> Accuracy: 0.5234375, Loss: 1.0668752193450928
Epoch: 0, Batch: 167 ==> Accuracy: 0.59375, Loss: 0.9804419279098511
Epoch: 0, Batch: 168 ==> Accuracy: 0.5390625, Loss: 1.0409290790557861
Epoch: 0, Batch: 169 ==> Accuracy: 0.5625, Loss: 1.0528626441955566
Epoch: 0, Batch: 170 ==> Accuracy: 0.544921875, Loss: 1.0455235242843628
Epoch: 0, Batch: 171 ==> Accuracy: 0.576171875, Loss: 1.0149381160736084
Epoch: 0, Batch: 172 ==> Accuracy: 0.576171875, Loss: 1.0148637294769287
Epoch: 0, Batch: 173 ==> Accuracy: 0.564453125, Loss: 1.0217645168304443
Epoch: 0, Batch: 174 ==> Accuracy: 0.55859375, Loss: 1.009751796722412
Epoch: 0, Batch: 175 ==> Accuracy: 0.546875, Loss: 1.0474194288253784
Epoch: 0, Batch: 176 ==> Accuracy: 0.587890625, Loss: 1.0148637294769287
Epoch: 0, Batch: 177 ==> Accuracy: 0.583984375, Loss: 1.029353141784668
Epoch: 0, Batch: 178 ==> Accuracy: 0.537109375, Loss: 1.0483585596084595
Epoch: 0, Batch: 179 ==> Accuracy: 0.53515625, Loss: 1.0493104457855225
Epoch: 0, Batch: 180 ==> Accuracy: 0.537109375, Loss: 1.0518248081207275
Epoch: 0, Batch: 181 ==> Accuracy: 0.529296875, Loss: 1.0885801315307617
Epoch: 0, Batch: 182 ==> Accuracy: 0.5546875, Loss: 1.0592478513717651
Epoch: 0, Batch: 183 ==> Accuracy: 0.580078125, Loss: 1.0194512605667114
Epoch: 0, Batch: 184 ==> Accuracy: 0.54296875, Loss: 1.0533742904663086
Epoch: 0, Batch: 185 ==> Accuracy: 0.572265625, Loss: 1.0364384651184082
Epoch: 0, Batch: 186 ==> Accuracy: 0.591796875, Loss: 1.0119330883026123
```

```
Epoch: 0, Batch: 187 ==> Accuracy: 0.560546875, Loss: 1.0353339910507202
Epoch: 0, Batch: 188 ==> Accuracy: 0.580078125, Loss: 0.9983108043670654
Epoch: 0, Batch: 189 ==> Accuracy: 0.560546875, Loss: 1.0466164350509644
Epoch: 0, Batch: 190 ==> Accuracy: 0.583984375, Loss: 0.9639889001846313
Epoch: 0, Batch: 191 ==> Accuracy: 0.591796875, Loss: 0.9911902546882629
Epoch: 0, Batch: 192 ==> Accuracy: 0.546875, Loss: 1.0295815467834473
Epoch: 0, Batch: 193 ==> Accuracy: 0.5546875, Loss: 1.0719525814056396
Epoch: 0, Batch: 194 ==> Accuracy: 0.548828125, Loss: 1.0401074886322021
Epoch: 0, Batch: 195 ==> Accuracy: 0.552734375, Loss: 1.0773365497589111
Epoch: 0, Batch: 196 ==> Accuracy: 0.55859375, Loss: 1.0449386835098267
Epoch: 0, Batch: 197 ==> Accuracy: 0.6015625, Loss: 0.9833877086639404
Epoch: 0, Batch: 198 ==> Accuracy: 0.5546875, Loss: 1.0330429077148438
Epoch: 0, Batch: 199 ==> Accuracy: 0.603515625, Loss: 0.9741150736808777
Epoch: 0, Batch: 200 ==> Accuracy: 0.55078125, Loss: 1.068802833557129
Epoch: 0, Batch: 201 ==> Accuracy: 0.5390625, Loss: 1.0761797428131104
Epoch: 0, Batch: 202 ==> Accuracy: 0.5859375, Loss: 1.0331716537475586
Epoch: 0, Batch: 203 ==> Accuracy: 0.564453125, Loss: 1.0015676021575928
Epoch: 0, Batch: 204 ==> Accuracy: 0.580078125, Loss: 0.9871879816055298
Epoch: 0, Batch: 205 ==> Accuracy: 0.578125, Loss: 1.0151913166046143
Epoch: 0, Batch: 206 ==> Accuracy: 0.5703125, Loss: 1.0135780572891235
Epoch: 0, Batch: 207 ==> Accuracy: 0.57421875, Loss: 0.9928833246231079
Epoch: 0, Batch: 208 ==> Accuracy: 0.55859375, Loss: 1.026843786239624
Epoch: 0, Batch: 209 ==> Accuracy: 0.615234375, Loss: 0.9722902774810791
Epoch: 0, Batch: 210 ==> Accuracy: 0.544921875, Loss: 1.0318903923034668
Epoch: 0, Batch: 211 ==> Accuracy: 0.58203125, Loss: 1.0057889223098755
Epoch: 0, Batch: 212 ==> Accuracy: 0.60546875, Loss: 0.9472935199737549
Epoch: 0, Batch: 213 ==> Accuracy: 0.541015625, Loss: 1.0716891288757324
Epoch: 0, Batch: 214 ==> Accuracy: 0.541015625, Loss: 1.0479483604431152
Epoch: 0, Batch: 215 ==> Accuracy: 0.546875, Loss: 1.0320160388946533
Epoch: 0, Batch: 216 ==> Accuracy: 0.5078125, Loss: 1.102677822113037
Epoch: 0, Batch: 217 ==> Accuracy: 0.580078125, Loss: 1.0153474807739258
Epoch: 0, Batch: 218 ==> Accuracy: 0.583984375, Loss: 1.0372443199157715
Epoch: 0, Batch: 219 ==> Accuracy: 0.568359375, Loss: 1.028740406036377
Epoch: 0, Batch: 220 ==> Accuracy: 0.53515625, Loss: 1.0980267524719238
Epoch: 0, Batch: 221 ==> Accuracy: 0.560546875, Loss: 1.0240956544876099
Epoch: 0, Batch: 222 ==> Accuracy: 0.580078125, Loss: 1.0533947944641113
Epoch: 0, Batch: 223 ==> Accuracy: 0.5703125, Loss: 0.9928874373435974
Epoch: 0, Batch: 224 ==> Accuracy: 0.578125, Loss: 1.018426775932312
Epoch: 0, Batch: 225 ==> Accuracy: 0.568359375, Loss: 1.0389580726623535
Epoch: 0, Batch: 226 ==> Accuracy: 0.6015625, Loss: 0.979727029800415
Epoch: 0, Batch: 227 ==> Accuracy: 0.53515625, Loss: 1.0477244853973389
Epoch: 0, Batch: 228 ==> Accuracy: 0.56640625, Loss: 1.0413771867752075
Epoch: 0, Batch: 229 ==> Accuracy: 0.55078125, Loss: 1.0462646484375
Epoch: 0, Batch: 230 ==> Accuracy: 0.59375, Loss: 1.0120341777801514
Epoch: 0, Batch: 231 ==> Accuracy: 0.583984375, Loss: 1.0413503646850586
Epoch: 0, Batch: 232 ==> Accuracy: 0.576171875, Loss: 0.9851844310760498
Epoch: 0, Batch: 233 ==> Accuracy: 0.603515625, Loss: 0.9889366626739502
Epoch: 0, Batch: 234 ==> Accuracy: 0.568359375, Loss: 1.0632226467132568
Epoch: 0, Batch: 235 ==> Accuracy: 0.591796875, Loss: 0.9551628828048706
Epoch: 0, Batch: 236 ==> Accuracy: 0.546875, Loss: 1.0509085655212402
Epoch: 0, Batch: 237 ==> Accuracy: 0.607421875, Loss: 0.986243486404419
Epoch: 0, Batch: 238 ==> Accuracy: 0.5546875, Loss: 1.0497815608978271
Epoch: 0, Batch: 239 ==> Accuracy: 0.5703125, Loss: 1.0200985670089722
Epoch: 0, Batch: 240 ==> Accuracy: 0.587890625, Loss: 1.0102183818817139
Epoch: 0, Batch: 241 ==> Accuracy: 0.615234375, Loss: 0.9635480642318726
```

```
Epoch: 0, Batch: 242 ==> Accuracy: 0.59765625, Loss: 0.9573809504508972
Epoch: 0, Batch: 243 ==> Accuracy: 0.6015625, Loss: 1.0130231380462646
Epoch: 0, Batch: 244 ==> Accuracy: 0.583984375, Loss: 0.987471342086792
Epoch: 0, Batch: 245 ==> Accuracy: 0.556640625, Loss: 1.0250811576843262
Epoch: 0, Batch: 246 ==> Accuracy: 0.5859375, Loss: 0.947037935256958
Epoch: 0, Batch: 247 ==> Accuracy: 0.60546875, Loss: 0.9448392391204834
Epoch: 0, Batch: 248 ==> Accuracy: 0.57421875, Loss: 1.0074127912521362
Epoch: 0, Batch: 249 ==> Accuracy: 0.580078125, Loss: 0.9947550892829895
Epoch: 0, Batch: 250 ==> Accuracy: 0.578125, Loss: 0.9955784678459167
Epoch: 0, Batch: 251 ==> Accuracy: 0.5625, Loss: 1.012412667274475
Epoch: 0, Batch: 252 ==> Accuracy: 0.58984375, Loss: 0.9848029017448425
Epoch: 0, Batch: 253 ==> Accuracy: 0.55078125, Loss: 1.0624016523361206
Epoch: 0, Batch: 254 ==> Accuracy: 0.5625, Loss: 1.0084044933319092
Epoch: 0, Batch: 255 ==> Accuracy: 0.490234375, Loss: 1.1507568359375
Epoch: 0, Batch: 256 ==> Accuracy: 0.576171875, Loss: 1.0347923040390015
Epoch: 0, Batch: 257 ==> Accuracy: 0.580078125, Loss: 1.0875539779663086
Epoch: 0, Batch: 258 ==> Accuracy: 0.552734375, Loss: 1.0572125911712646
Epoch: 0, Batch: 259 ==> Accuracy: 0.572265625, Loss: 1.0059103965759277
Epoch: 0, Batch: 260 ==> Accuracy: 0.564453125, Loss: 1.0638353824615479
Epoch: 0, Validation ==> Accuracy: 0.564453125, Loss: 1.0638353824615479
Epoch: 1, Batch: 0 ==> Accuracy: 0.59375, Loss: 0.9821502566337585
Epoch: 1, Batch: 1 ==> Accuracy: 0.51953125, Loss: 1.0419849157333374
Epoch: 1, Batch: 2 ==> Accuracy: 0.6015625, Loss: 1.0047860145568848
Epoch: 1, Batch: 3 ==> Accuracy: 0.57421875, Loss: 1.0176379680633545
Epoch: 1, Batch: 4 ==> Accuracy: 0.5859375, Loss: 0.9840190410614014
Epoch: 1, Batch: 5 ==> Accuracy: 0.58203125, Loss: 1.009450912475586
Epoch: 1, Batch: 6 ==> Accuracy: 0.5703125, Loss: 1.0393290519714355
Epoch: 1, Batch: 7 ==> Accuracy: 0.59375, Loss: 1.0124480724334717
Epoch: 1, Batch: 8 ==> Accuracy: 0.591796875, Loss: 0.9736231565475464
Epoch: 1, Batch: 9 ==> Accuracy: 0.583984375, Loss: 0.9845097064971924
Epoch: 1, Batch: 10 ==> Accuracy: 0.564453125, Loss: 0.9956216812133789
Epoch: 1, Batch: 11 ==> Accuracy: 0.59765625, Loss: 0.9871234893798828
Epoch: 1, Batch: 12 ==> Accuracy: 0.59375, Loss: 0.9635047912597656
Epoch: 1, Batch: 13 ==> Accuracy: 0.583984375, Loss: 1.0016791820526123
Epoch: 1, Batch: 14 ==> Accuracy: 0.59765625, Loss: 0.9709572792053223
Epoch: 1, Batch: 15 ==> Accuracy: 0.57421875, Loss: 0.9778182506561279
Epoch: 1, Batch: 16 ==> Accuracy: 0.58203125, Loss: 0.9667856693267822
Epoch: 1, Batch: 17 ==> Accuracy: 0.5859375, Loss: 0.9775816202163696
Epoch: 1, Batch: 18 ==> Accuracy: 0.623046875, Loss: 0.928144097328186
Epoch: 1, Batch: 19 ==> Accuracy: 0.599609375, Loss: 1.0115325450897217
Epoch: 1, Batch: 20 ==> Accuracy: 0.59765625, Loss: 0.9608590602874756
Epoch: 1, Batch: 21 ==> Accuracy: 0.576171875, Loss: 1.0303324460983276
Epoch: 1, Batch: 22 ==> Accuracy: 0.56640625, Loss: 1.0169507265090942
Epoch: 1, Batch: 23 ==> Accuracy: 0.591796875, Loss: 1.0195895433425903
Epoch: 1, Batch: 24 ==> Accuracy: 0.591796875, Loss: 0.9586073756217957
Epoch: 1, Batch: 25 ==> Accuracy: 0.626953125, Loss: 0.8971009254455566
Epoch: 1, Batch: 26 ==> Accuracy: 0.609375, Loss: 0.9518213272094727
Epoch: 1, Batch: 27 ==> Accuracy: 0.568359375, Loss: 0.9962042570114136
Epoch: 1, Batch: 28 ==> Accuracy: 0.58203125, Loss: 1.000859260559082
Epoch: 1, Batch: 29 ==> Accuracy: 0.56640625, Loss: 1.0147686004638672
Epoch: 1, Batch: 30 ==> Accuracy: 0.5859375, Loss: 1.0245298147201538
Epoch: 1, Batch: 31 ==> Accuracy: 0.5703125, Loss: 0.9832244515419006
Epoch: 1, Batch: 32 ==> Accuracy: 0.6015625, Loss: 0.9447551965713501
Epoch: 1, Batch: 33 ==> Accuracy: 0.5859375, Loss: 0.9750351309776306
Epoch: 1, Batch: 34 ==> Accuracy: 0.572265625, Loss: 0.994066596031189
```

```
Epoch: 1, Batch: 35 ==> Accuracy: 0.60546875, Loss: 0.9330116510391235
Epoch: 1, Batch: 36 ==> Accuracy: 0.623046875, Loss: 0.938536524772644
Epoch: 1, Batch: 37 ==> Accuracy: 0.595703125, Loss: 0.9889663457870483
Epoch: 1, Batch: 38 ==> Accuracy: 0.59375, Loss: 0.9766460657119751
Epoch: 1, Batch: 39 ==> Accuracy: 0.583984375, Loss: 0.9864119291305542
Epoch: 1, Batch: 40 ==> Accuracy: 0.62109375, Loss: 0.927419126033783
Epoch: 1, Batch: 41 ==> Accuracy: 0.595703125, Loss: 0.9983661770820618
Epoch: 1, Batch: 42 ==> Accuracy: 0.583984375, Loss: 0.979023814201355
Epoch: 1, Batch: 43 ==> Accuracy: 0.587890625, Loss: 0.967949628829956
Epoch: 1, Batch: 44 ==> Accuracy: 0.578125, Loss: 0.9959185123443604
Epoch: 1, Batch: 45 ==> Accuracy: 0.5625, Loss: 1.0160623788833618
Epoch: 1, Batch: 46 ==> Accuracy: 0.57421875, Loss: 0.9993640780448914
Epoch: 1, Batch: 47 ==> Accuracy: 0.544921875, Loss: 1.0277870893478394
Epoch: 1, Batch: 48 ==> Accuracy: 0.5859375, Loss: 0.9934065341949463
Epoch: 1, Batch: 49 ==> Accuracy: 0.583984375, Loss: 0.9976356029510498
Epoch: 1, Batch: 50 ==> Accuracy: 0.5625, Loss: 1.004457712173462
Epoch: 1, Batch: 51 ==> Accuracy: 0.611328125, Loss: 0.9374815225601196
Epoch: 1, Batch: 52 ==> Accuracy: 0.55859375, Loss: 1.0220894813537598
Epoch: 1, Batch: 53 ==> Accuracy: 0.623046875, Loss: 0.9235490560531616
Epoch: 1, Batch: 54 ==> Accuracy: 0.55859375, Loss: 1.0035321712493896
Epoch: 1, Batch: 55 ==> Accuracy: 0.587890625, Loss: 0.9724566340446472
Epoch: 1, Batch: 56 ==> Accuracy: 0.609375, Loss: 0.9469468593597412
Epoch: 1, Batch: 57 ==> Accuracy: 0.603515625, Loss: 0.9589359760284424
Epoch: 1, Batch: 58 ==> Accuracy: 0.599609375, Loss: 0.9776930809020996
Epoch: 1, Batch: 59 ==> Accuracy: 0.56640625, Loss: 0.9929418563842773
Epoch: 1, Batch: 60 ==> Accuracy: 0.595703125, Loss: 0.9623795747756958
Epoch: 1, Batch: 61 ==> Accuracy: 0.58203125, Loss: 0.9914960265159607
Epoch: 1, Batch: 62 ==> Accuracy: 0.58984375, Loss: 0.9645367860794067
Epoch: 1, Batch: 63 ==> Accuracy: 0.609375, Loss: 0.9504888653755188
Epoch: 1, Batch: 64 ==> Accuracy: 0.599609375, Loss: 0.9628129005432129
Epoch: 1, Batch: 65 ==> Accuracy: 0.578125, Loss: 0.9960894584655762
Epoch: 1, Batch: 66 ==> Accuracy: 0.6015625, Loss: 0.9651522040367126
Epoch: 1, Batch: 67 ==> Accuracy: 0.595703125, Loss: 0.9111549854278564
Epoch: 1, Batch: 68 ==> Accuracy: 0.611328125, Loss: 0.9564224481582642
Epoch: 1, Batch: 69 ==> Accuracy: 0.58203125, Loss: 0.9886537790298462
Epoch: 1, Batch: 70 ==> Accuracy: 0.626953125, Loss: 0.9015594720840454
Epoch: 1, Batch: 71 ==> Accuracy: 0.642578125, Loss: 0.9191182851791382
Epoch: 1, Batch: 72 ==> Accuracy: 0.544921875, Loss: 1.0134313106536865
Epoch: 1, Batch: 73 ==> Accuracy: 0.59375, Loss: 1.0088587999343872
Epoch: 1, Batch: 74 ==> Accuracy: 0.57421875, Loss: 1.0010490417480469
Epoch: 1, Batch: 75 ==> Accuracy: 0.580078125, Loss: 0.9555443525314331
Epoch: 1, Batch: 76 ==> Accuracy: 0.56640625, Loss: 1.0334490537643433
Epoch: 1, Batch: 77 ==> Accuracy: 0.564453125, Loss: 1.0291911363601685
Epoch: 1, Batch: 78 ==> Accuracy: 0.583984375, Loss: 0.9961630702018738
Epoch: 1, Batch: 79 ==> Accuracy: 0.595703125, Loss: 0.9941927790641785
Epoch: 1, Batch: 80 ==> Accuracy: 0.58203125, Loss: 0.9876503348350525
Epoch: 1, Batch: 81 ==> Accuracy: 0.572265625, Loss: 1.001116156578064
Epoch: 1, Batch: 82 ==> Accuracy: 0.572265625, Loss: 1.0413165092468262
Epoch: 1, Batch: 83 ==> Accuracy: 0.599609375, Loss: 0.9782854318618774
Epoch: 1, Batch: 84 ==> Accuracy: 0.619140625, Loss: 0.9542613625526428
Epoch: 1, Batch: 85 ==> Accuracy: 0.54296875, Loss: 1.0181667804718018
Epoch: 1, Batch: 86 ==> Accuracy: 0.6015625, Loss: 0.9559391736984253
Epoch: 1, Batch: 87 ==> Accuracy: 0.5859375, Loss: 0.979465901851654
Epoch: 1, Batch: 88 ==> Accuracy: 0.62890625, Loss: 0.9545197486877441
Epoch: 1, Batch: 89 ==> Accuracy: 0.56640625, Loss: 0.995428204536438
```

```
Epoch: 1, Batch: 90 ==> Accuracy: 0.6015625, Loss: 0.9855406880378723
Epoch: 1, Batch: 91 ==> Accuracy: 0.591796875, Loss: 1.0319325923919678
Epoch: 1, Batch: 92 ==> Accuracy: 0.5703125, Loss: 1.0096192359924316
Epoch: 1, Batch: 93 ==> Accuracy: 0.55078125, Loss: 1.0331993103027344
Epoch: 1, Batch: 94 ==> Accuracy: 0.5859375, Loss: 0.949618935585022
Epoch: 1, Batch: 95 ==> Accuracy: 0.578125, Loss: 0.9968045949935913
Epoch: 1, Batch: 96 ==> Accuracy: 0.611328125, Loss: 0.9324296712875366
Epoch: 1, Batch: 97 ==> Accuracy: 0.57421875, Loss: 0.9831117391586304
Epoch: 1, Batch: 98 ==> Accuracy: 0.611328125, Loss: 0.9463702440261841
Epoch: 1, Batch: 99 ==> Accuracy: 0.6015625, Loss: 0.9706481695175171
Epoch: 1, Batch: 100 ==> Accuracy: 0.5859375, Loss: 1.0171592235565186
Epoch: 1, Batch: 101 ==> Accuracy: 0.580078125, Loss: 0.9703892469406128
Epoch: 1, Batch: 102 ==> Accuracy: 0.6328125, Loss: 0.900244951248169
Epoch: 1, Batch: 103 ==> Accuracy: 0.59765625, Loss: 0.9762716889381409
Epoch: 1, Batch: 104 ==> Accuracy: 0.638671875, Loss: 0.906129002571106
Epoch: 1, Batch: 105 ==> Accuracy: 0.611328125, Loss: 0.9310550689697266
Epoch: 1, Batch: 106 ==> Accuracy: 0.607421875, Loss: 0.9816278219223022
Epoch: 1, Batch: 107 ==> Accuracy: 0.619140625, Loss: 0.8994443416595459
Epoch: 1, Batch: 108 ==> Accuracy: 0.603515625, Loss: 0.9746057391166687
Epoch: 1, Batch: 109 ==> Accuracy: 0.619140625, Loss: 0.950664758682251
Epoch: 1, Batch: 110 ==> Accuracy: 0.6171875, Loss: 0.9331191778182983
Epoch: 1, Batch: 111 ==> Accuracy: 0.5703125, Loss: 1.0045690536499023
Epoch: 1, Batch: 112 ==> Accuracy: 0.5703125, Loss: 0.9938790798187256
Epoch: 1, Batch: 113 ==> Accuracy: 0.556640625, Loss: 1.0296002626419067
Epoch: 1, Batch: 114 ==> Accuracy: 0.59375, Loss: 0.97258460521698
Epoch: 1, Batch: 115 ==> Accuracy: 0.5859375, Loss: 1.00941801071167
Epoch: 1, Batch: 116 ==> Accuracy: 0.580078125, Loss: 0.9577515125274658
Epoch: 1, Batch: 117 ==> Accuracy: 0.5625, Loss: 1.0051743984222412
Epoch: 1, Batch: 118 ==> Accuracy: 0.576171875, Loss: 1.0146751403808594
Epoch: 1, Batch: 119 ==> Accuracy: 0.603515625, Loss: 0.9456072449684143
Epoch: 1, Batch: 120 ==> Accuracy: 0.630859375, Loss: 0.9380902051925659
Epoch: 1, Batch: 121 ==> Accuracy: 0.640625, Loss: 0.9076797366142273
Epoch: 1, Batch: 122 ==> Accuracy: 0.6015625, Loss: 0.9998031854629517
Epoch: 1, Batch: 123 ==> Accuracy: 0.60546875, Loss: 0.9359471797943115
Epoch: 1, Batch: 124 ==> Accuracy: 0.58984375, Loss: 0.959995448589325
Epoch: 1, Batch: 125 ==> Accuracy: 0.59375, Loss: 0.9733254909515381
Epoch: 1, Batch: 126 ==> Accuracy: 0.60546875, Loss: 0.95814448595047
Epoch: 1, Batch: 127 ==> Accuracy: 0.603515625, Loss: 0.9748998284339905
Epoch: 1, Batch: 128 ==> Accuracy: 0.62109375, Loss: 0.9084258079528809
Epoch: 1, Batch: 129 ==> Accuracy: 0.609375, Loss: 0.9618493318557739
Epoch: 1, Batch: 130 ==> Accuracy: 0.609375, Loss: 0.9766512513160706
Epoch: 1, Batch: 131 ==> Accuracy: 0.61328125, Loss: 0.954799473285675
Epoch: 1, Batch: 132 ==> Accuracy: 0.599609375, Loss: 0.9480630159378052
Epoch: 1, Batch: 133 ==> Accuracy: 0.578125, Loss: 0.9596871733665466
Epoch: 1, Batch: 134 ==> Accuracy: 0.60546875, Loss: 0.9316904544830322
Epoch: 1, Batch: 135 ==> Accuracy: 0.623046875, Loss: 0.9076110124588013
Epoch: 1, Batch: 136 ==> Accuracy: 0.60546875, Loss: 0.9649897217750549
Epoch: 1, Batch: 137 ==> Accuracy: 0.62890625, Loss: 0.9451881647109985
Epoch: 1, Batch: 138 ==> Accuracy: 0.630859375, Loss: 0.9048124551773071
Epoch: 1, Batch: 139 ==> Accuracy: 0.630859375, Loss: 0.9179521799087524
Epoch: 1, Batch: 140 ==> Accuracy: 0.595703125, Loss: 0.9458411931991577
Epoch: 1, Batch: 141 ==> Accuracy: 0.6015625, Loss: 0.942423939704895
Epoch: 1, Batch: 142 ==> Accuracy: 0.55859375, Loss: 1.002780795097351
Epoch: 1, Batch: 143 ==> Accuracy: 0.607421875, Loss: 0.958351850596436
Epoch: 1, Batch: 144 ==> Accuracy: 0.640625, Loss: 0.9010056853294373
```

```
Epoch: 1, Batch: 145 ==> Accuracy: 0.572265625, Loss: 1.005915641784668
Epoch: 1, Batch: 146 ==> Accuracy: 0.61328125, Loss: 0.9378282427787781
Epoch: 1, Batch: 147 ==> Accuracy: 0.634765625, Loss: 0.8991029262542725
Epoch: 1, Batch: 148 ==> Accuracy: 0.583984375, Loss: 0.9247543811798096
Epoch: 1, Batch: 149 ==> Accuracy: 0.609375, Loss: 0.9128595590591431
Epoch: 1, Batch: 150 ==> Accuracy: 0.625, Loss: 0.9350258708000183
Epoch: 1, Batch: 151 ==> Accuracy: 0.607421875, Loss: 0.9676188230514526
Epoch: 1, Batch: 152 ==> Accuracy: 0.58203125, Loss: 0.9330538511276245
Epoch: 1, Batch: 153 ==> Accuracy: 0.59375, Loss: 0.9545934200286865
Epoch: 1, Batch: 154 ==> Accuracy: 0.634765625, Loss: 0.9005956649780273
Epoch: 1, Batch: 155 ==> Accuracy: 0.578125, Loss: 0.9819316864013672
Epoch: 1, Batch: 156 ==> Accuracy: 0.55859375, Loss: 1.0326225757598877
Epoch: 1, Batch: 157 ==> Accuracy: 0.638671875, Loss: 0.8743561506271362
Epoch: 1, Batch: 158 ==> Accuracy: 0.58203125, Loss: 0.9795804023742676
Epoch: 1, Batch: 159 ==> Accuracy: 0.63671875, Loss: 0.9217538833618164
Epoch: 1, Batch: 160 ==> Accuracy: 0.68359375, Loss: 0.8511516451835632
Epoch: 1, Batch: 161 ==> Accuracy: 0.591796875, Loss: 0.9521121978759766
Epoch: 1, Batch: 162 ==> Accuracy: 0.6171875, Loss: 0.9313716888427734
Epoch: 1, Batch: 163 ==> Accuracy: 0.6171875, Loss: 0.9135153889656067
Epoch: 1, Batch: 164 ==> Accuracy: 0.61328125, Loss: 0.9221986532211304
Epoch: 1, Batch: 165 ==> Accuracy: 0.638671875, Loss: 0.9078529477119446
Epoch: 1, Batch: 166 ==> Accuracy: 0.578125, Loss: 0.9673219919204712
Epoch: 1, Batch: 167 ==> Accuracy: 0.6171875, Loss: 0.9232171177864075
Epoch: 1, Batch: 168 ==> Accuracy: 0.578125, Loss: 1.0303256511688232
Epoch: 1, Batch: 169 ==> Accuracy: 0.556640625, Loss: 1.031794548034668
Epoch: 1, Batch: 170 ==> Accuracy: 0.61328125, Loss: 0.9356296062469482
Epoch: 1, Batch: 171 ==> Accuracy: 0.6171875, Loss: 0.96979820728302
Epoch: 1, Batch: 172 ==> Accuracy: 0.6328125, Loss: 0.9213032722473145
Epoch: 1, Batch: 173 ==> Accuracy: 0.6171875, Loss: 0.936229944229126
Epoch: 1, Batch: 174 ==> Accuracy: 0.59375, Loss: 0.965034544467926
Epoch: 1, Batch: 175 ==> Accuracy: 0.6171875, Loss: 0.9164213538169861
Epoch: 1, Batch: 176 ==> Accuracy: 0.603515625, Loss: 0.9410723447799683
Epoch: 1, Batch: 177 ==> Accuracy: 0.615234375, Loss: 0.9178826808929443
Epoch: 1, Batch: 178 ==> Accuracy: 0.56640625, Loss: 0.9931444525718689
Epoch: 1, Batch: 179 ==> Accuracy: 0.60546875, Loss: 0.9522731304168701
Epoch: 1, Batch: 180 ==> Accuracy: 0.60546875, Loss: 0.9470561742782593
Epoch: 1, Batch: 181 ==> Accuracy: 0.615234375, Loss: 0.9164619445800781
Epoch: 1, Batch: 182 ==> Accuracy: 0.57421875, Loss: 0.9910396933555603
Epoch: 1, Batch: 183 ==> Accuracy: 0.611328125, Loss: 0.9482900500297546
Epoch: 1, Batch: 184 ==> Accuracy: 0.619140625, Loss: 0.9708278179168701
Epoch: 1, Batch: 185 ==> Accuracy: 0.587890625, Loss: 0.9456285834312439
Epoch: 1, Batch: 186 ==> Accuracy: 0.626953125, Loss: 0.9063573479652405
Epoch: 1, Batch: 187 ==> Accuracy: 0.626953125, Loss: 0.905576229095459
Epoch: 1, Batch: 188 ==> Accuracy: 0.591796875, Loss: 0.9597083330154419
Epoch: 1, Batch: 189 ==> Accuracy: 0.5859375, Loss: 0.9378684759140015
Epoch: 1, Batch: 190 ==> Accuracy: 0.60546875, Loss: 0.9181869626045227
Epoch: 1, Batch: 191 ==> Accuracy: 0.619140625, Loss: 0.9412652850151062
Epoch: 1, Batch: 192 ==> Accuracy: 0.62890625, Loss: 0.9392780065536499
Epoch: 1, Batch: 193 ==> Accuracy: 0.583984375, Loss: 0.9829510450363159
Epoch: 1, Batch: 194 ==> Accuracy: 0.619140625, Loss: 0.9531497955322266
Epoch: 1, Batch: 195 ==> Accuracy: 0.59765625, Loss: 0.9560445547103882
Epoch: 1, Batch: 196 ==> Accuracy: 0.59375, Loss: 0.9865173697471619
Epoch: 1, Batch: 197 ==> Accuracy: 0.638671875, Loss: 0.8848099708557129
Epoch: 1, Batch: 198 ==> Accuracy: 0.591796875, Loss: 0.9768030047416687
Epoch: 1, Batch: 199 ==> Accuracy: 0.60546875, Loss: 0.9068001508712769
```

```
Epoch: 1, Batch: 200 ==> Accuracy: 0.6171875, Loss: 0.9422774910926819
Epoch: 1, Batch: 201 ==> Accuracy: 0.591796875, Loss: 0.9793338775634766
Epoch: 1, Batch: 202 ==> Accuracy: 0.603515625, Loss: 0.9196930527687073
Epoch: 1, Batch: 203 ==> Accuracy: 0.6171875, Loss: 0.94843590259552
Epoch: 1, Batch: 204 ==> Accuracy: 0.615234375, Loss: 0.9343578815460205
Epoch: 1, Batch: 205 ==> Accuracy: 0.5859375, Loss: 0.9568079710006714
Epoch: 1, Batch: 206 ==> Accuracy: 0.591796875, Loss: 0.9550272226333618
Epoch: 1, Batch: 207 ==> Accuracy: 0.630859375, Loss: 0.8874225616455078
Epoch: 1, Batch: 208 ==> Accuracy: 0.599609375, Loss: 0.9575220346450806
Epoch: 1, Batch: 209 ==> Accuracy: 0.611328125, Loss: 0.9526613354682922
Epoch: 1, Batch: 210 ==> Accuracy: 0.595703125, Loss: 0.9254511594772339
Epoch: 1, Batch: 211 ==> Accuracy: 0.623046875, Loss: 0.8990371823310852
Epoch: 1, Batch: 212 ==> Accuracy: 0.61328125, Loss: 0.8980031609535217
Epoch: 1, Batch: 213 ==> Accuracy: 0.564453125, Loss: 0.9720073938369751
Epoch: 1, Batch: 214 ==> Accuracy: 0.626953125, Loss: 0.9291844367980957
Epoch: 1, Batch: 215 ==> Accuracy: 0.611328125, Loss: 0.972628116607666
Epoch: 1, Batch: 216 ==> Accuracy: 0.599609375, Loss: 0.9603784084320068
Epoch: 1, Batch: 217 ==> Accuracy: 0.638671875, Loss: 0.9254639148712158
Epoch: 1, Batch: 218 ==> Accuracy: 0.6171875, Loss: 0.9288874864578247
Epoch: 1, Batch: 219 ==> Accuracy: 0.6171875, Loss: 0.9250667095184326
Epoch: 1, Batch: 220 ==> Accuracy: 0.615234375, Loss: 0.9460748434066772
Epoch: 1, Batch: 221 ==> Accuracy: 0.646484375, Loss: 0.8699826598167419
Epoch: 1, Batch: 222 ==> Accuracy: 0.62109375, Loss: 0.9268887042999268
Epoch: 1, Batch: 223 ==> Accuracy: 0.63671875, Loss: 0.9190798997879028
Epoch: 1, Batch: 224 ==> Accuracy: 0.59765625, Loss: 0.9810799360275269
Epoch: 1, Batch: 225 ==> Accuracy: 0.640625, Loss: 0.8889449834823608
Epoch: 1, Batch: 226 ==> Accuracy: 0.58203125, Loss: 1.0141197443008423
Epoch: 1, Batch: 227 ==> Accuracy: 0.59375, Loss: 0.9901441335678101
Epoch: 1, Batch: 228 ==> Accuracy: 0.62890625, Loss: 0.9017726182937622
Epoch: 1, Batch: 229 ==> Accuracy: 0.603515625, Loss: 0.9222848415374756
Epoch: 1, Batch: 230 ==> Accuracy: 0.63671875, Loss: 0.9118143916130066
Epoch: 1, Batch: 231 ==> Accuracy: 0.568359375, Loss: 1.007826328277588
Epoch: 1, Batch: 232 ==> Accuracy: 0.59765625, Loss: 0.9678477644920349
Epoch: 1, Batch: 233 ==> Accuracy: 0.61328125, Loss: 0.9436293840408325
Epoch: 1, Batch: 234 ==> Accuracy: 0.60546875, Loss: 0.9128298759460449
Epoch: 1, Batch: 235 ==> Accuracy: 0.615234375, Loss: 0.940295934677124
Epoch: 1, Batch: 236 ==> Accuracy: 0.62890625, Loss: 0.898622989654541
Epoch: 1, Batch: 237 ==> Accuracy: 0.59765625, Loss: 0.9564652442932129
Epoch: 1, Batch: 238 ==> Accuracy: 0.623046875, Loss: 0.9368937015533447
Epoch: 1, Batch: 239 ==> Accuracy: 0.564453125, Loss: 1.0528361797332764
Epoch: 1, Batch: 240 ==> Accuracy: 0.63671875, Loss: 0.9131319522857666
Epoch: 1, Batch: 241 ==> Accuracy: 0.6171875, Loss: 0.9415488243103027
Epoch: 1, Batch: 242 ==> Accuracy: 0.619140625, Loss: 0.8883740901947021
Epoch: 1, Batch: 243 ==> Accuracy: 0.619140625, Loss: 0.9335421323776245
Epoch: 1, Batch: 244 ==> Accuracy: 0.59765625, Loss: 0.9359438419342041
Epoch: 1, Batch: 245 ==> Accuracy: 0.599609375, Loss: 0.9253177642822266
Epoch: 1, Batch: 246 ==> Accuracy: 0.591796875, Loss: 0.9481754302978516
Epoch: 1, Batch: 247 ==> Accuracy: 0.626953125, Loss: 0.8926298022270203
Epoch: 1, Batch: 248 ==> Accuracy: 0.6171875, Loss: 0.9458264112472534
Epoch: 1, Batch: 249 ==> Accuracy: 0.609375, Loss: 0.9124596118927002
Epoch: 1, Batch: 250 ==> Accuracy: 0.626953125, Loss: 0.897932231426239
Epoch: 1, Batch: 251 ==> Accuracy: 0.6171875, Loss: 0.9006710052490234
Epoch: 1, Batch: 252 ==> Accuracy: 0.583984375, Loss: 0.9526329040527344
Epoch: 1, Batch: 253 ==> Accuracy: 0.609375, Loss: 0.9247385263442993
Epoch: 1, Batch: 254 ==> Accuracy: 0.619140625, Loss: 0.8689848184585571
```

```
Epoch: 1, Batch: 255 ==> Accuracy: 0.603515625, Loss: 0.946711003780365
Epoch: 1, Batch: 256 ==> Accuracy: 0.58984375, Loss: 0.991351842880249
Epoch: 1, Batch: 257 ==> Accuracy: 0.61328125, Loss: 0.8895043134689331
Epoch: 1, Batch: 258 ==> Accuracy: 0.623046875, Loss: 0.9064812064170837
Epoch: 1, Batch: 259 ==> Accuracy: 0.61328125, Loss: 0.9169475436210632
Epoch: 1, Batch: 260 ==> Accuracy: 0.623046875, Loss: 0.930535078048706
Epoch: 1, Validation ==> Accuracy: 0.623046875, Loss: 0.930535078048706
Epoch: 2, Batch: 0 ==> Accuracy: 0.640625, Loss: 0.8738363981246948
Epoch: 2, Batch: 1 ==> Accuracy: 0.666015625, Loss: 0.876366138458252
Epoch: 2, Batch: 2 ==> Accuracy: 0.62890625, Loss: 0.928436815738678
Epoch: 2, Batch: 3 ==> Accuracy: 0.61328125, Loss: 0.9171689748764038
Epoch: 2, Batch: 4 ==> Accuracy: 0.607421875, Loss: 0.9363100528717041
Epoch: 2, Batch: 5 ==> Accuracy: 0.625, Loss: 0.9061827063560486
Epoch: 2, Batch: 6 ==> Accuracy: 0.62109375, Loss: 0.9425342082977295
Epoch: 2, Batch: 7 ==> Accuracy: 0.595703125, Loss: 0.949606716632843
Epoch: 2, Batch: 8 ==> Accuracy: 0.60546875, Loss: 0.9307500720024109
Epoch: 2, Batch: 9 ==> Accuracy: 0.642578125, Loss: 0.9132318496704102
Epoch: 2, Batch: 10 ==> Accuracy: 0.607421875, Loss: 0.9478812217712402
Epoch: 2, Batch: 11 ==> Accuracy: 0.59375, Loss: 0.9128555059432983
Epoch: 2, Batch: 12 ==> Accuracy: 0.591796875, Loss: 0.9282547235488892
Epoch: 2, Batch: 13 ==> Accuracy: 0.6015625, Loss: 0.928859531879425
Epoch: 2, Batch: 14 ==> Accuracy: 0.62109375, Loss: 0.9270974397659302
Epoch: 2, Batch: 15 ==> Accuracy: 0.62890625, Loss: 0.9212660193443298
Epoch: 2, Batch: 16 ==> Accuracy: 0.646484375, Loss: 0.8720869421958923
Epoch: 2, Batch: 17 ==> Accuracy: 0.62109375, Loss: 0.9237233400344849
Epoch: 2, Batch: 18 ==> Accuracy: 0.59765625, Loss: 0.9548066854476929
Epoch: 2, Batch: 19 ==> Accuracy: 0.6328125, Loss: 0.8783406019210815
Epoch: 2, Batch: 20 ==> Accuracy: 0.59375, Loss: 0.9112462997436523
Epoch: 2, Batch: 21 ==> Accuracy: 0.58203125, Loss: 0.9507560729980469
Epoch: 2, Batch: 22 ==> Accuracy: 0.634765625, Loss: 0.9723023176193237
Epoch: 2, Batch: 23 ==> Accuracy: 0.609375, Loss: 0.9495068192481995
Epoch: 2, Batch: 24 ==> Accuracy: 0.5859375, Loss: 0.9324507713317871
Epoch: 2, Batch: 25 ==> Accuracy: 0.6015625, Loss: 0.9587597846984863
Epoch: 2, Batch: 26 ==> Accuracy: 0.611328125, Loss: 0.8971563577651978
Epoch: 2, Batch: 27 ==> Accuracy: 0.64453125, Loss: 0.8752658367156982
Epoch: 2, Batch: 28 ==> Accuracy: 0.65625, Loss: 0.856776237487793
Epoch: 2, Batch: 29 ==> Accuracy: 0.599609375, Loss: 0.9353039860725403
Epoch: 2, Batch: 30 ==> Accuracy: 0.6171875, Loss: 0.9232318997383118
Epoch: 2, Batch: 31 ==> Accuracy: 0.619140625, Loss: 0.9177266359329224
Epoch: 2, Batch: 32 ==> Accuracy: 0.599609375, Loss: 0.9454358220100403
Epoch: 2, Batch: 33 ==> Accuracy: 0.625, Loss: 0.9065042734146118
Epoch: 2, Batch: 34 ==> Accuracy: 0.654296875, Loss: 0.8948538899421692
Epoch: 2, Batch: 35 ==> Accuracy: 0.634765625, Loss: 0.8800727128982544
Epoch: 2, Batch: 36 ==> Accuracy: 0.623046875, Loss: 0.9016320109367371
Epoch: 2, Batch: 37 ==> Accuracy: 0.650390625, Loss: 0.8036226034164429
Epoch: 2, Batch: 38 ==> Accuracy: 0.630859375, Loss: 0.9089432954788208
Epoch: 2, Batch: 39 ==> Accuracy: 0.63671875, Loss: 0.875115156173706
Epoch: 2, Batch: 40 ==> Accuracy: 0.640625, Loss: 0.8828781843185425
Epoch: 2, Batch: 41 ==> Accuracy: 0.638671875, Loss: 0.9313891530036926
Epoch: 2, Batch: 42 ==> Accuracy: 0.615234375, Loss: 0.952964186668396
Epoch: 2, Batch: 43 ==> Accuracy: 0.623046875, Loss: 0.9397825598716736
Epoch: 2, Batch: 44 ==> Accuracy: 0.671875, Loss: 0.8502488732337952
Epoch: 2, Batch: 45 ==> Accuracy: 0.615234375, Loss: 0.9577433466911316
Epoch: 2, Batch: 46 ==> Accuracy: 0.650390625, Loss: 0.9060928821563721
Epoch: 2, Batch: 47 ==> Accuracy: 0.58984375, Loss: 0.9469389319419861
```

```
Epoch: 2, Batch: 48 ==> Accuracy: 0.59375, Loss: 0.9439125061035156
Epoch: 2, Batch: 49 ==> Accuracy: 0.630859375, Loss: 0.912418007850647
Epoch: 2, Batch: 50 ==> Accuracy: 0.6171875, Loss: 0.9135178327560425
Epoch: 2, Batch: 51 ==> Accuracy: 0.630859375, Loss: 0.8873916864395142
Epoch: 2, Batch: 52 ==> Accuracy: 0.642578125, Loss: 0.9107084274291992
Epoch: 2, Batch: 53 ==> Accuracy: 0.62890625, Loss: 0.9066836833953857
Epoch: 2, Batch: 54 ==> Accuracy: 0.615234375, Loss: 0.9328285455703735
Epoch: 2, Batch: 55 ==> Accuracy: 0.62890625, Loss: 0.9071819186210632
Epoch: 2, Batch: 56 ==> Accuracy: 0.642578125, Loss: 0.8630607724189758
Epoch: 2, Batch: 57 ==> Accuracy: 0.625, Loss: 0.939259946346283
Epoch: 2, Batch: 58 ==> Accuracy: 0.6328125, Loss: 0.8609874248504639
Epoch: 2, Batch: 59 ==> Accuracy: 0.634765625, Loss: 0.8853731155395508
Epoch: 2, Batch: 60 ==> Accuracy: 0.634765625, Loss: 0.8924157619476318
Epoch: 2, Batch: 61 ==> Accuracy: 0.615234375, Loss: 0.9224322438240051
Epoch: 2, Batch: 62 ==> Accuracy: 0.671875, Loss: 0.8301726579666138
Epoch: 2, Batch: 63 ==> Accuracy: 0.640625, Loss: 0.8296281099319458
Epoch: 2, Batch: 64 ==> Accuracy: 0.650390625, Loss: 0.8499056696891785
Epoch: 2, Batch: 65 ==> Accuracy: 0.61328125, Loss: 0.9239330887794495
Epoch: 2, Batch: 66 ==> Accuracy: 0.646484375, Loss: 0.8586957454681396
Epoch: 2, Batch: 67 ==> Accuracy: 0.615234375, Loss: 0.9130194187164307
Epoch: 2, Batch: 68 ==> Accuracy: 0.609375, Loss: 0.9589753150939941
Epoch: 2, Batch: 69 ==> Accuracy: 0.62109375, Loss: 0.9708173274993896
Epoch: 2, Batch: 70 ==> Accuracy: 0.625, Loss: 0.8842142820358276
Epoch: 2, Batch: 71 ==> Accuracy: 0.591796875, Loss: 0.9875202178955078
Epoch: 2, Batch: 72 ==> Accuracy: 0.599609375, Loss: 0.9633340835571289
Epoch: 2, Batch: 73 ==> Accuracy: 0.63671875, Loss: 0.9017676711082458
Epoch: 2, Batch: 74 ==> Accuracy: 0.640625, Loss: 0.8966745138168335
Epoch: 2, Batch: 75 ==> Accuracy: 0.630859375, Loss: 0.8917003273963928
Epoch: 2, Batch: 76 ==> Accuracy: 0.61328125, Loss: 0.8977159857749939
Epoch: 2, Batch: 77 ==> Accuracy: 0.68359375, Loss: 0.8245815634727478
Epoch: 2, Batch: 78 ==> Accuracy: 0.603515625, Loss: 0.9511623382568359
Epoch: 2, Batch: 79 ==> Accuracy: 0.615234375, Loss: 0.9437236189842224
Epoch: 2, Batch: 80 ==> Accuracy: 0.60546875, Loss: 0.9817422032356262
Epoch: 2, Batch: 81 ==> Accuracy: 0.6015625, Loss: 0.9088102579116821
Epoch: 2, Batch: 82 ==> Accuracy: 0.578125, Loss: 0.990249752998352
Epoch: 2, Batch: 83 ==> Accuracy: 0.63671875, Loss: 0.8485457301139832
Epoch: 2, Batch: 84 ==> Accuracy: 0.578125, Loss: 0.9718008041381836
Epoch: 2, Batch: 85 ==> Accuracy: 0.6171875, Loss: 0.9371365308761597
Epoch: 2, Batch: 86 ==> Accuracy: 0.560546875, Loss: 0.9937323927879333
Epoch: 2, Batch: 87 ==> Accuracy: 0.638671875, Loss: 0.9148505926132202
Epoch: 2, Batch: 88 ==> Accuracy: 0.609375, Loss: 0.965650200843811
Epoch: 2, Batch: 89 ==> Accuracy: 0.583984375, Loss: 0.9822380542755127
Epoch: 2, Batch: 90 ==> Accuracy: 0.640625, Loss: 0.8651560544967651
Epoch: 2, Batch: 91 ==> Accuracy: 0.60546875, Loss: 0.9080208539962769
Epoch: 2, Batch: 92 ==> Accuracy: 0.6171875, Loss: 0.9159311056137085
Epoch: 2, Batch: 93 ==> Accuracy: 0.65234375, Loss: 0.8649896383285522
Epoch: 2, Batch: 94 ==> Accuracy: 0.650390625, Loss: 0.8754371404647827
Epoch: 2, Batch: 95 ==> Accuracy: 0.642578125, Loss: 0.86558997631073
Epoch: 2, Batch: 96 ==> Accuracy: 0.625, Loss: 0.8944637775421143
Epoch: 2, Batch: 97 ==> Accuracy: 0.630859375, Loss: 0.8745276927947998
Epoch: 2, Batch: 98 ==> Accuracy: 0.640625, Loss: 0.921581506729126
Epoch: 2, Batch: 99 ==> Accuracy: 0.640625, Loss: 0.8995109796524048
Epoch: 2, Batch: 100 ==> Accuracy: 0.64453125, Loss: 0.8703320026397705
Epoch: 2, Batch: 101 ==> Accuracy: 0.623046875, Loss: 0.9461137652397156
Epoch: 2, Batch: 102 ==> Accuracy: 0.6640625, Loss: 0.8812047839164734
```

```
Epoch: 2, Batch: 103 ==> Accuracy: 0.619140625, Loss: 0.9099265336990356
Epoch: 2, Batch: 104 ==> Accuracy: 0.595703125, Loss: 0.9820626974105835
Epoch: 2, Batch: 105 ==> Accuracy: 0.658203125, Loss: 0.8668367862701416
Epoch: 2, Batch: 106 ==> Accuracy: 0.662109375, Loss: 0.8703039288520813
Epoch: 2, Batch: 107 ==> Accuracy: 0.64453125, Loss: 0.8874421119689941
Epoch: 2, Batch: 108 ==> Accuracy: 0.634765625, Loss: 0.90590500831604
Epoch: 2, Batch: 109 ==> Accuracy: 0.640625, Loss: 0.8365345597267151
Epoch: 2, Batch: 110 ==> Accuracy: 0.59765625, Loss: 0.9566037654876709
Epoch: 2, Batch: 111 ==> Accuracy: 0.599609375, Loss: 0.9585572481155396
Epoch: 2, Batch: 112 ==> Accuracy: 0.630859375, Loss: 0.9096163511276245
Epoch: 2, Batch: 113 ==> Accuracy: 0.62109375, Loss: 0.9504913091659546
Epoch: 2, Batch: 114 ==> Accuracy: 0.61328125, Loss: 0.8883659839630127
Epoch: 2, Batch: 115 ==> Accuracy: 0.580078125, Loss: 0.9910285472869873
Epoch: 2, Batch: 116 ==> Accuracy: 0.61328125, Loss: 0.9617281556129456
Epoch: 2, Batch: 117 ==> Accuracy: 0.6015625, Loss: 0.9250712394714355
Epoch: 2, Batch: 118 ==> Accuracy: 0.6328125, Loss: 0.8770548701286316
Epoch: 2, Batch: 119 ==> Accuracy: 0.62109375, Loss: 0.9341816306114197
Epoch: 2, Batch: 120 ==> Accuracy: 0.625, Loss: 0.9262001514434814
Epoch: 2, Batch: 121 ==> Accuracy: 0.611328125, Loss: 0.8911091685295105
Epoch: 2, Batch: 122 ==> Accuracy: 0.62109375, Loss: 0.9151989221572876
Epoch: 2, Batch: 123 ==> Accuracy: 0.62109375, Loss: 0.969207227230072
Epoch: 2, Batch: 124 ==> Accuracy: 0.611328125, Loss: 0.9136056900024414
Epoch: 2, Batch: 125 ==> Accuracy: 0.603515625, Loss: 0.9322187304496765
Epoch: 2, Batch: 126 ==> Accuracy: 0.646484375, Loss: 0.9135769009590149
Epoch: 2, Batch: 127 ==> Accuracy: 0.60546875, Loss: 0.9214986562728882
Epoch: 2, Batch: 128 ==> Accuracy: 0.625, Loss: 0.9234912395477295
Epoch: 2, Batch: 129 ==> Accuracy: 0.662109375, Loss: 0.8824090957641602
Epoch: 2, Batch: 130 ==> Accuracy: 0.63671875, Loss: 0.9166477918624878
Epoch: 2, Batch: 131 ==> Accuracy: 0.634765625, Loss: 0.9047625064849854
Epoch: 2, Batch: 132 ==> Accuracy: 0.634765625, Loss: 0.8671693801879883
Epoch: 2, Batch: 133 ==> Accuracy: 0.6328125, Loss: 0.876706600189209
Epoch: 2, Batch: 134 ==> Accuracy: 0.580078125, Loss: 0.9429401159286499
Epoch: 2, Batch: 135 ==> Accuracy: 0.67578125, Loss: 0.850955605506897
Epoch: 2, Batch: 136 ==> Accuracy: 0.640625, Loss: 0.8546554446220398
Epoch: 2, Batch: 137 ==> Accuracy: 0.58203125, Loss: 0.9351463317871094
Epoch: 2, Batch: 138 ==> Accuracy: 0.626953125, Loss: 0.8627094626426697
Epoch: 2, Batch: 139 ==> Accuracy: 0.625, Loss: 0.9166244864463806
Epoch: 2, Batch: 140 ==> Accuracy: 0.626953125, Loss: 0.8941432237625122
Epoch: 2, Batch: 141 ==> Accuracy: 0.623046875, Loss: 0.8930957317352295
Epoch: 2, Batch: 142 ==> Accuracy: 0.630859375, Loss: 0.9119506478309631
Epoch: 2, Batch: 143 ==> Accuracy: 0.63671875, Loss: 0.9065364599227905
Epoch: 2, Batch: 144 ==> Accuracy: 0.638671875, Loss: 0.8796790838241577
Epoch: 2, Batch: 145 ==> Accuracy: 0.634765625, Loss: 0.9311784505844116
Epoch: 2, Batch: 146 ==> Accuracy: 0.61328125, Loss: 0.9089457988739014
Epoch: 2, Batch: 147 ==> Accuracy: 0.609375, Loss: 0.9172720909118652
Epoch: 2, Batch: 148 ==> Accuracy: 0.650390625, Loss: 0.8798774480819702
Epoch: 2, Batch: 149 ==> Accuracy: 0.58984375, Loss: 0.9211869239807129
Epoch: 2, Batch: 150 ==> Accuracy: 0.6328125, Loss: 0.9200294613838196
Epoch: 2, Batch: 151 ==> Accuracy: 0.6171875, Loss: 0.9222480058670044
Epoch: 2, Batch: 152 ==> Accuracy: 0.630859375, Loss: 0.869739294052124
Epoch: 2, Batch: 153 ==> Accuracy: 0.654296875, Loss: 0.8690295219421387
Epoch: 2, Batch: 154 ==> Accuracy: 0.640625, Loss: 0.8858442902565002
Epoch: 2, Batch: 155 ==> Accuracy: 0.685546875, Loss: 0.8295618891716003
Epoch: 2, Batch: 156 ==> Accuracy: 0.677734375, Loss: 0.8592317700386047
Epoch: 2, Batch: 157 ==> Accuracy: 0.642578125, Loss: 0.8738887310028076
```

```
Epoch: 2, Batch: 158 ==> Accuracy: 0.61328125, Loss: 0.8826789855957031
Epoch: 2, Batch: 159 ==> Accuracy: 0.626953125, Loss: 0.9105467796325684
Epoch: 2, Batch: 160 ==> Accuracy: 0.625, Loss: 0.8909931182861328
Epoch: 2, Batch: 161 ==> Accuracy: 0.65234375, Loss: 0.8337885141372681
Epoch: 2, Batch: 162 ==> Accuracy: 0.62109375, Loss: 0.9291839599609375
Epoch: 2, Batch: 163 ==> Accuracy: 0.646484375, Loss: 0.8396455645561218
Epoch: 2, Batch: 164 ==> Accuracy: 0.638671875, Loss: 0.8576846122741699
Epoch: 2, Batch: 165 ==> Accuracy: 0.6796875, Loss: 0.8226604461669922
Epoch: 2, Batch: 166 ==> Accuracy: 0.62890625, Loss: 0.8678005933761597
Epoch: 2, Batch: 167 ==> Accuracy: 0.6171875, Loss: 0.9023828506469727
Epoch: 2, Batch: 168 ==> Accuracy: 0.66796875, Loss: 0.8378435969352722
Epoch: 2, Batch: 169 ==> Accuracy: 0.62890625, Loss: 0.9375811219215393
Epoch: 2, Batch: 170 ==> Accuracy: 0.654296875, Loss: 0.8643175363540649
Epoch: 2, Batch: 171 ==> Accuracy: 0.65625, Loss: 0.8501081466674805
Epoch: 2, Batch: 172 ==> Accuracy: 0.673828125, Loss: 0.8537993431091309
Epoch: 2, Batch: 173 ==> Accuracy: 0.6171875, Loss: 0.8889597654342651
Epoch: 2, Batch: 174 ==> Accuracy: 0.58203125, Loss: 0.9809814691543579
Epoch: 2, Batch: 175 ==> Accuracy: 0.634765625, Loss: 0.8841351270675659
Epoch: 2, Batch: 176 ==> Accuracy: 0.650390625, Loss: 0.8443059921264648
Epoch: 2, Batch: 177 ==> Accuracy: 0.630859375, Loss: 0.8668565154075623
Epoch: 2, Batch: 178 ==> Accuracy: 0.662109375, Loss: 0.8667718768119812
Epoch: 2, Batch: 179 ==> Accuracy: 0.599609375, Loss: 0.9433473348617554
Epoch: 2, Batch: 180 ==> Accuracy: 0.626953125, Loss: 0.8750351667404175
Epoch: 2, Batch: 181 ==> Accuracy: 0.640625, Loss: 0.8782007694244385
Epoch: 2, Batch: 182 ==> Accuracy: 0.609375, Loss: 0.9363893270492554
Epoch: 2, Batch: 183 ==> Accuracy: 0.6328125, Loss: 0.8685743808746338
Epoch: 2, Batch: 184 ==> Accuracy: 0.634765625, Loss: 0.8955647945404053
Epoch: 2, Batch: 185 ==> Accuracy: 0.640625, Loss: 0.9228205680847168
Epoch: 2, Batch: 186 ==> Accuracy: 0.64453125, Loss: 0.8510221242904663
Epoch: 2, Batch: 187 ==> Accuracy: 0.615234375, Loss: 0.93611741065979
Epoch: 2, Batch: 188 ==> Accuracy: 0.638671875, Loss: 0.8950217962265015
Epoch: 2, Batch: 189 ==> Accuracy: 0.642578125, Loss: 0.865541934967041
Epoch: 2, Batch: 190 ==> Accuracy: 0.65625, Loss: 0.8729898929595947
Epoch: 2, Batch: 191 ==> Accuracy: 0.642578125, Loss: 0.8592928647994995
Epoch: 2, Batch: 192 ==> Accuracy: 0.654296875, Loss: 0.8520129919052124
Epoch: 2, Batch: 193 ==> Accuracy: 0.666015625, Loss: 0.8411188125610352
Epoch: 2, Batch: 194 ==> Accuracy: 0.62109375, Loss: 0.8709500432014465
Epoch: 2, Batch: 195 ==> Accuracy: 0.62109375, Loss: 0.9100310802459717
Epoch: 2, Batch: 196 ==> Accuracy: 0.65625, Loss: 0.862663745880127
Epoch: 2, Batch: 197 ==> Accuracy: 0.6171875, Loss: 0.9165470004081726
Epoch: 2, Batch: 198 ==> Accuracy: 0.65625, Loss: 0.8596616983413696
Epoch: 2, Batch: 199 ==> Accuracy: 0.634765625, Loss: 0.8453389406204224
Epoch: 2, Batch: 200 ==> Accuracy: 0.642578125, Loss: 0.8669220209121704
Epoch: 2, Batch: 201 ==> Accuracy: 0.583984375, Loss: 0.9606139659881592
Epoch: 2, Batch: 202 ==> Accuracy: 0.6484375, Loss: 0.8319884538650513
Epoch: 2, Batch: 203 ==> Accuracy: 0.64453125, Loss: 0.8356696963310242
Epoch: 2, Batch: 204 ==> Accuracy: 0.580078125, Loss: 0.9559433460235596
Epoch: 2, Batch: 205 ==> Accuracy: 0.646484375, Loss: 0.8762895464897156
Epoch: 2, Batch: 206 ==> Accuracy: 0.634765625, Loss: 0.8356398344039917
Epoch: 2, Batch: 207 ==> Accuracy: 0.626953125, Loss: 0.9114456176757812
Epoch: 2, Batch: 208 ==> Accuracy: 0.669921875, Loss: 0.7836244106292725
Epoch: 2, Batch: 209 ==> Accuracy: 0.642578125, Loss: 0.8439240455627441
Epoch: 2, Batch: 210 ==> Accuracy: 0.634765625, Loss: 0.8645399212837219
Epoch: 2, Batch: 211 ==> Accuracy: 0.658203125, Loss: 0.8554019927978516
Epoch: 2, Batch: 212 ==> Accuracy: 0.642578125, Loss: 0.8784623146057129
```

```
Epoch: 2, Batch: 213 ==> Accuracy: 0.658203125, Loss: 0.835344672203064
Epoch: 2, Batch: 214 ==> Accuracy: 0.68359375, Loss: 0.8638163208961487
Epoch: 2, Batch: 215 ==> Accuracy: 0.607421875, Loss: 0.9224132299423218
Epoch: 2, Batch: 216 ==> Accuracy: 0.662109375, Loss: 0.8594542741775513
Epoch: 2, Batch: 217 ==> Accuracy: 0.62890625, Loss: 0.9067776203155518
Epoch: 2, Batch: 218 ==> Accuracy: 0.662109375, Loss: 0.8340948224067688
Epoch: 2, Batch: 219 ==> Accuracy: 0.681640625, Loss: 0.8217395544052124
Epoch: 2, Batch: 220 ==> Accuracy: 0.615234375, Loss: 0.878846287727356
Epoch: 2, Batch: 221 ==> Accuracy: 0.6328125, Loss: 0.8798737525939941
Epoch: 2, Batch: 222 ==> Accuracy: 0.66015625, Loss: 0.8742822408676147
Epoch: 2, Batch: 223 ==> Accuracy: 0.640625, Loss: 0.8871294260025024
Epoch: 2, Batch: 224 ==> Accuracy: 0.61328125, Loss: 0.9028563499450684
Epoch: 2, Batch: 225 ==> Accuracy: 0.6171875, Loss: 0.8970466256141663
Epoch: 2, Batch: 226 ==> Accuracy: 0.6171875, Loss: 0.8804556131362915
Epoch: 2, Batch: 227 ==> Accuracy: 0.62109375, Loss: 0.9010173082351685
Epoch: 2, Batch: 228 ==> Accuracy: 0.640625, Loss: 0.8730192184448242
Epoch: 2, Batch: 229 ==> Accuracy: 0.6171875, Loss: 0.9055677652359009
Epoch: 2, Batch: 230 ==> Accuracy: 0.646484375, Loss: 0.8588941097259521
Epoch: 2, Batch: 231 ==> Accuracy: 0.650390625, Loss: 0.8934353590011597
Epoch: 2, Batch: 232 ==> Accuracy: 0.69140625, Loss: 0.7844000458717346
Epoch: 2, Batch: 233 ==> Accuracy: 0.66796875, Loss: 0.8332664370536804
Epoch: 2, Batch: 234 ==> Accuracy: 0.623046875, Loss: 0.8788613080978394
Epoch: 2, Batch: 235 ==> Accuracy: 0.6328125, Loss: 0.8710574507713318
Epoch: 2, Batch: 236 ==> Accuracy: 0.666015625, Loss: 0.829903781414032
Epoch: 2, Batch: 237 ==> Accuracy: 0.66015625, Loss: 0.8075175285339355
Epoch: 2, Batch: 238 ==> Accuracy: 0.681640625, Loss: 0.8245964646339417
Epoch: 2, Batch: 239 ==> Accuracy: 0.625, Loss: 0.8882259130477905
Epoch: 2, Batch: 240 ==> Accuracy: 0.64453125, Loss: 0.8802674412727356
Epoch: 2, Batch: 241 ==> Accuracy: 0.6328125, Loss: 0.8930946588516235
Epoch: 2, Batch: 242 ==> Accuracy: 0.62890625, Loss: 0.91059410572052
Epoch: 2, Batch: 243 ==> Accuracy: 0.646484375, Loss: 0.8943105936050415
Epoch: 2, Batch: 244 ==> Accuracy: 0.666015625, Loss: 0.7992074489593506
Epoch: 2, Batch: 245 ==> Accuracy: 0.666015625, Loss: 0.8341249823570251
Epoch: 2, Batch: 246 ==> Accuracy: 0.640625, Loss: 0.8663008809089661
Epoch: 2, Batch: 247 ==> Accuracy: 0.599609375, Loss: 0.9199360609054565
Epoch: 2, Batch: 248 ==> Accuracy: 0.666015625, Loss: 0.8151282072067261
Epoch: 2, Batch: 249 ==> Accuracy: 0.65234375, Loss: 0.8635216951370239
Epoch: 2, Batch: 250 ==> Accuracy: 0.638671875, Loss: 0.9085956811904907
Epoch: 2, Batch: 251 ==> Accuracy: 0.63671875, Loss: 0.8687394857406616
Epoch: 2, Batch: 252 ==> Accuracy: 0.642578125, Loss: 0.8840270638465881
Epoch: 2, Batch: 253 ==> Accuracy: 0.63671875, Loss: 0.8869205713272095
Epoch: 2, Batch: 254 ==> Accuracy: 0.626953125, Loss: 0.8701528906822205
Epoch: 2, Batch: 255 ==> Accuracy: 0.66796875, Loss: 0.7972009181976318
Epoch: 2, Batch: 256 ==> Accuracy: 0.650390625, Loss: 0.8959128856658936
Epoch: 2, Batch: 257 ==> Accuracy: 0.64453125, Loss: 0.8569722771644592
Epoch: 2, Batch: 258 ==> Accuracy: 0.640625, Loss: 0.825919508934021
Epoch: 2, Batch: 259 ==> Accuracy: 0.66796875, Loss: 0.8939763307571411
Epoch: 2, Batch: 260 ==> Accuracy: 0.638671875, Loss: 0.9527095556259155
Epoch: 2, Validation ==> Accuracy: 0.638671875, Loss: 0.9527095556259155
Epoch: 3, Batch: 0 ==> Accuracy: 0.66015625, Loss: 0.8567299842834473
Epoch: 3, Batch: 1 ==> Accuracy: 0.63671875, Loss: 0.8346589803695679
Epoch: 3, Batch: 2 ==> Accuracy: 0.662109375, Loss: 0.8376191854476929
Epoch: 3, Batch: 3 ==> Accuracy: 0.669921875, Loss: 0.8140338659286499
Epoch: 3, Batch: 4 ==> Accuracy: 0.650390625, Loss: 0.850626528263092
Epoch: 3, Batch: 5 ==> Accuracy: 0.642578125, Loss: 0.8916327953338623
```

```
Epoch: 3, Batch: 6 ==> Accuracy: 0.67578125, Loss: 0.7879724502563477
Epoch: 3, Batch: 7 ==> Accuracy: 0.69921875, Loss: 0.801372766494751
Epoch: 3, Batch: 8 ==> Accuracy: 0.701171875, Loss: 0.7762398719787598
Epoch: 3, Batch: 9 ==> Accuracy: 0.654296875, Loss: 0.8562706708908081
Epoch: 3, Batch: 10 ==> Accuracy: 0.65234375, Loss: 0.8079544305801392
Epoch: 3, Batch: 11 ==> Accuracy: 0.650390625, Loss: 0.8088281154632568
Epoch: 3, Batch: 12 ==> Accuracy: 0.61328125, Loss: 0.9160823822021484
Epoch: 3, Batch: 13 ==> Accuracy: 0.630859375, Loss: 0.8880718350410461
Epoch: 3, Batch: 14 ==> Accuracy: 0.640625, Loss: 0.8633613586425781
Epoch: 3, Batch: 15 ==> Accuracy: 0.65625, Loss: 0.83968186378479
Epoch: 3, Batch: 16 ==> Accuracy: 0.669921875, Loss: 0.7866041660308838
Epoch: 3, Batch: 17 ==> Accuracy: 0.63671875, Loss: 0.8599061965942383
Epoch: 3, Batch: 18 ==> Accuracy: 0.630859375, Loss: 0.85429847240448
Epoch: 3, Batch: 19 ==> Accuracy: 0.66796875, Loss: 0.8508899211883545
Epoch: 3, Batch: 20 ==> Accuracy: 0.68359375, Loss: 0.814898669719696
Epoch: 3, Batch: 21 ==> Accuracy: 0.6328125, Loss: 0.876140296459198
Epoch: 3, Batch: 22 ==> Accuracy: 0.6484375, Loss: 0.8571212291717529
Epoch: 3, Batch: 23 ==> Accuracy: 0.681640625, Loss: 0.7954908609390259
Epoch: 3, Batch: 24 ==> Accuracy: 0.67578125, Loss: 0.8112902045249939
Epoch: 3, Batch: 25 ==> Accuracy: 0.650390625, Loss: 0.8360112905502319
Epoch: 3, Batch: 26 ==> Accuracy: 0.681640625, Loss: 0.7602719068527222
Epoch: 3, Batch: 27 ==> Accuracy: 0.638671875, Loss: 0.8640182018280029
Epoch: 3, Batch: 28 ==> Accuracy: 0.646484375, Loss: 0.8520387411117554
Epoch: 3, Batch: 29 ==> Accuracy: 0.630859375, Loss: 0.888916015625
Epoch: 3, Batch: 30 ==> Accuracy: 0.6484375, Loss: 0.8422170877456665
Epoch: 3, Batch: 31 ==> Accuracy: 0.671875, Loss: 0.8081763982772827
Epoch: 3, Batch: 32 ==> Accuracy: 0.634765625, Loss: 0.8971860408782959
Epoch: 3, Batch: 33 ==> Accuracy: 0.66796875, Loss: 0.8070653676986694
Epoch: 3, Batch: 34 ==> Accuracy: 0.689453125, Loss: 0.7678910493850708
Epoch: 3, Batch: 35 ==> Accuracy: 0.650390625, Loss: 0.8718198537826538
```